

Lab Assignment-4

Q.1 Implement `fuzzy_c_means` ($X, c, m=2.0, \text{max_iter}=100, \text{tol}=1e-5$) from scratch in Python (do NOT use external libraries like `fcmeans`).

Algorithm steps:

1. Initialize a random membership matrix $U \in \mathbb{R}^{N \times c}$, ensuring each row sums to 1.
2. Update cluster centers V using the weighted average data points.
3. Update membership values based on distances and fuzzifier m .
4. Stop when the maximum change in U is less than `tol` or `max_iter` is reached.
5. Return the final cluster centers, the membership matrix, and the number of iterations used. Test the function on a 2-D blobs dataset (`sklearn.datasets.make_blobs`, 3 clusters, 300 points) and plot:
 - The data points colored by their highest membership
 - The Cluster centers

Q.2 Apply your `fuzzy_c_means()` implementation to a grayscale image segmentation task.

Steps:

1. Load a grayscale image (e.g., from Pillow) and flatten it to a 1-D array.
2. Apply FCM with 3 clusters.
3. Assign each pixel to the cluster with highest membership and reconstruct the segmented image.
4. Display the original and segmented images side-by-side.
5. Report the runtime and mean squared difference between original and segmented image.

Q.3 Investigate the effect of fuzzifier m on clustering.

For a fixed dataset (`make_blobs` with 3 clusters):

1. Run your FCM for $m = 1.5, 2.0, 2.5, 3.0$.
2. For each, report:
 - Final objective function value
 - Number of iterations
 - Plot of data colored by dominant membership

Write a short analysis: how does increasing m affect membership sharpness and convergence?

Q.4 Use Partition Coefficient (PC) and Partition Entropy (PE) as validity indices.

Implement:

1. $\text{partition_coefficient}(U) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^c u_{ij}^2$
2. $\text{partition_entropy}(U) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^c u_{ij} \log u_{ij}$

For $k = 2 \dots 6$, run FCM on the Iris dataset (`load_iris()`) and plot PC and PE vs. k on a single plot (one curve increasing, one decreasing).
Select and report the best k based on both indices.

Q.5 Perform document clustering with FCM using TF-IDF vectors.

Steps:

1. Load up to 3,000 documents from `sklearn.datasets.fetch_20newsgroups` (select "rec.autos", "sci.space", "talk.politics.misc").
2. Compute TF-IDF vectors (`TfidfVectorizer`) and apply dimensionality reduction (`TruncatedSVD`, 50 components).
3. Run FCM with $c = 3$.
4. Assign cluster labels by highest membership and compare against true categories using Adjusted Rand Index (ARI).
5. Report ARI and display a bar chart of average membership per cluster.